

# Exercise 5 R Markdown

Santanu Mukherjee

07/26/2022

## R Markdown

### Question

This question should be answered using the “*Weekly*” data set, which is part of the “**ISLR**” package. This data is similar in nature to the “*Smarket*” data from this chapter’s lab, except that it contains 1089 weekly returns for 21 years, from the beginning of 1990 to the end of 2010.

#### Question a

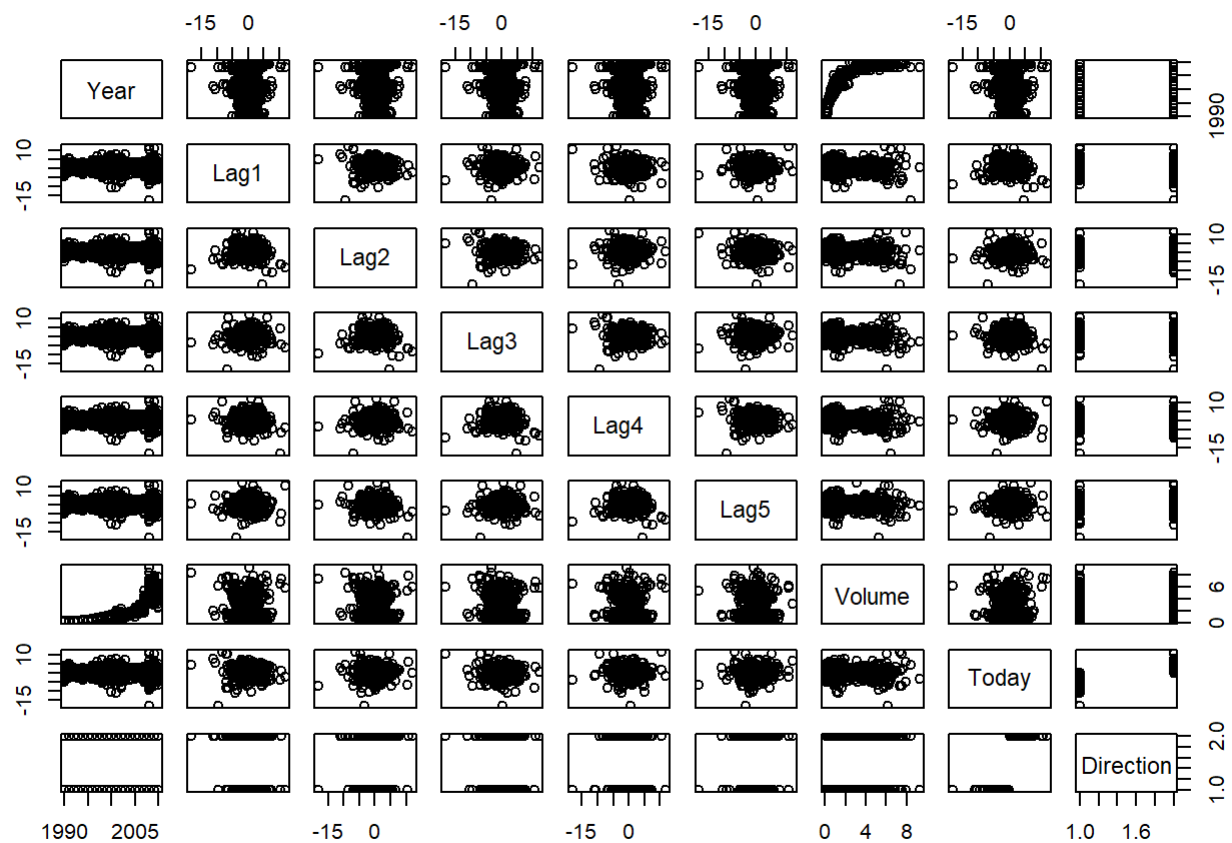
- a. Produce some numerical and graphical summaries of the “*Weekly*” data. Do there appear to be any patterns ?

### Answer a

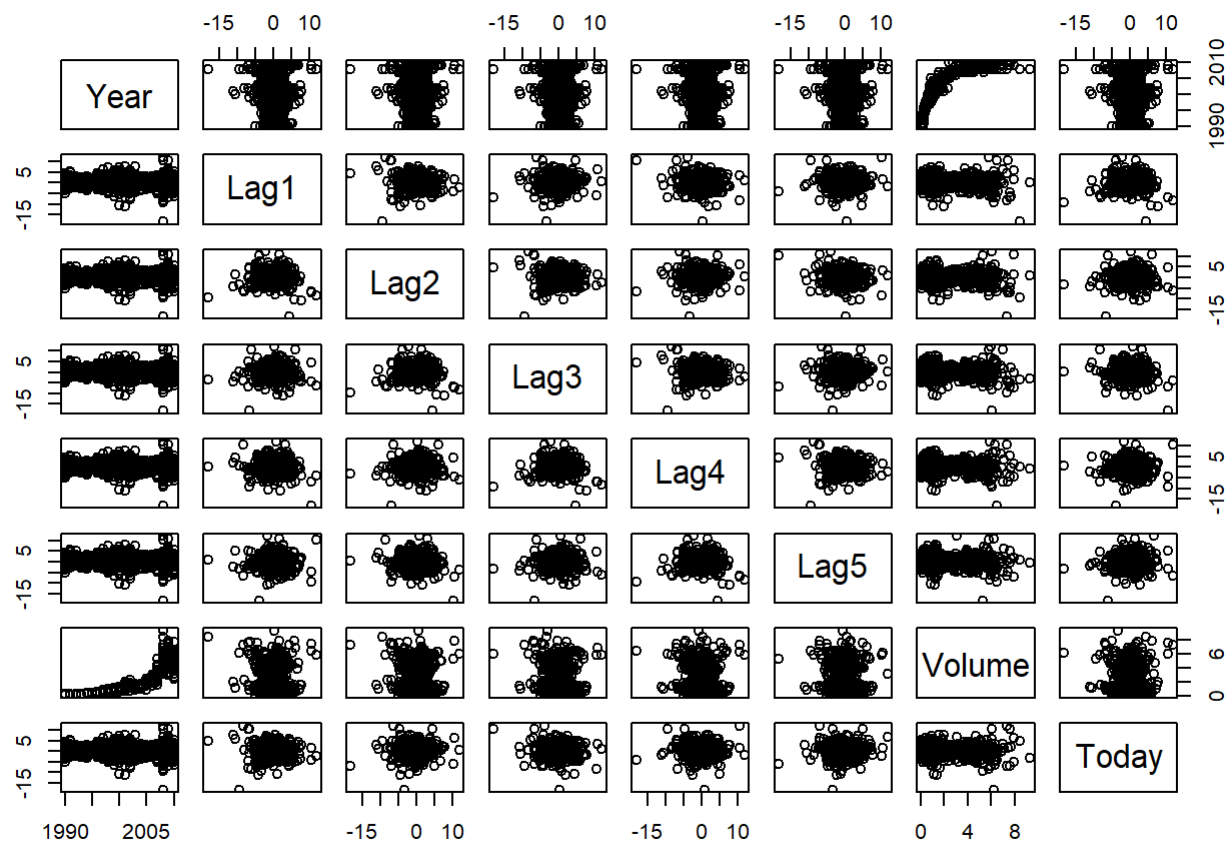
```
##Part (a) Weekly Data Summary  
library(ISLR)  
summary(Weekly)
```

```
##      Year      Lag1      Lag2      Lag3
## Min.   :1990  Min.   :-18.1950  Min.   :-18.1950  Min.   :-18.1950
## 1st Qu.:1995  1st Qu.: -1.1540  1st Qu.: -1.1540  1st Qu.: -1.1580
## Median :2000  Median :  0.2410  Median :  0.2410  Median :  0.2410
## Mean   :2000  Mean   :  0.1506  Mean   :  0.1511  Mean   :  0.1472
## 3rd Qu.:2005  3rd Qu.:  1.4050  3rd Qu.:  1.4090  3rd Qu.:  1.4090
## Max.   :2010  Max.   : 12.0260  Max.   : 12.0260  Max.   : 12.0260
##      Lag4      Lag5      Volume      Today
## Min.   :-18.1950  Min.   :-18.1950  Min.   :0.08747  Min.   :-18.1950
## 1st Qu.: -1.1580  1st Qu.: -1.1660  1st Qu.:0.33202  1st Qu.: -1.1540
## Median :  0.2380  Median :  0.2340  Median :1.00268  Median :  0.2410
## Mean   :  0.1458  Mean   :  0.1399  Mean   :1.57462  Mean   :  0.1499
## 3rd Qu.:  1.4090  3rd Qu.:  1.4050  3rd Qu.:2.05373  3rd Qu.:  1.4050
## Max.   : 12.0260  Max.   : 12.0260  Max.   :9.32821  Max.   : 12.0260
## Direction
## Down:484
## Up :605
##
##
##
##
```

```
pairs(Weekly)
```




```
pairs(Weekly[, -9])
```



```
cor(Weekly[, -9])
```

```
##           Year      Lag1      Lag2      Lag3      Lag4
## Year      1.00000000 -0.032289274 -0.03339001 -0.03000649 -0.031127923
## Lag1     -0.03228927  1.000000000 -0.07485305  0.05863568 -0.071273876
## Lag2     -0.03339001 -0.074853051  1.00000000 -0.07572091  0.058381535
## Lag3     -0.03000649  0.058635682 -0.07572091  1.00000000 -0.075395865
## Lag4     -0.03112792 -0.071273876  0.05838153 -0.07539587  1.000000000
## Lag5     -0.03051910 -0.008183096 -0.07249948  0.06065717 -0.075675027
## Volume    0.84194162 -0.064951313 -0.08551314 -0.06928771 -0.061074617
## Today    -0.03245989 -0.075031842  0.05916672 -0.07124364 -0.007825873
##           Lag5      Volume      Today
## Year    -0.030519101  0.84194162 -0.032459894
## Lag1    -0.008183096 -0.06495131 -0.075031842
## Lag2    -0.072499482 -0.08551314  0.059166717
## Lag3     0.060657175 -0.06928771 -0.071243639
## Lag4    -0.075675027 -0.06107462 -0.007825873
## Lag5     1.000000000 -0.05851741  0.011012698
## Volume  -0.058517414  1.00000000 -0.033077783
## Today    0.011012698 -0.03307778  1.000000000
```

```
attach(Weekly)
plot(Volume, col="green")
```



Step by Step Observations:

1. The *Summary* and subsequently the *pairs* showed that the variable “Direction” was insignificant.
2. So, then I got the correlation matrix with all variables except **Direction**.
3. The correlations between the “**lag**” variables and **Today\* variable are close to zero**.
- 4.\*\* The correlation between variables “**Year**” and “**Volume**” is the only significant one.
5. So, I have done plot “Volume”, and I see that is increasing over time.

## Question b

- b. Use the full data set to perform a logistic regression with “*Direction*” as the response and the five lag variables plus “*Volume*” as predictors.  
Use the summary function to print the results. Do any of the predictors appear to be statistically significant ? If so, which ones ?

## Answer b

### ##Part (b) Logistic Regression

```
set.seed(1)
log.reg <- glm(Direction~Lag1+Lag2+Lag3+Lag4+Lag5+Volume, data=Weekly,family=binomial)
summary(log.reg)
```

```
##
## Call:
## glm(formula = Direction ~ Lag1 + Lag2 + Lag3 + Lag4 + Lag5 +
##      Volume, family = binomial, data = Weekly)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -1.6949  -1.2565   0.9913   1.0849   1.4579
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)  0.26686    0.08593   3.106  0.0019 **
## Lag1        -0.04127    0.02641  -1.563   0.1181
## Lag2         0.05844    0.02686   2.175   0.0296 *
## Lag3        -0.01606    0.02666  -0.602   0.5469
## Lag4        -0.02779    0.02646  -1.050   0.2937
## Lag5        -0.01447    0.02638  -0.549   0.5833
## Volume       -0.02274    0.03690  -0.616   0.5377
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 1496.2  on 1088  degrees of freedom
## Residual deviance: 1486.4  on 1082  degrees of freedom
## AIC: 1500.4
##
## Number of Fisher Scoring iterations: 4
```

It seems that “Lag2” is the only predictor which is statistically significant at  $\alpha = 0.05$  as its p-value is less than 0.05.

## Question c

- c. Compute the confusion matrix and overall fraction of correct predictions. Explain what the confusion matrix is telling you about the types of mistakes made by logistic regression.

## Answer c

```
##Part (c) Confusion Matrix
```

```
prob.log.reg <- predict(log.reg, type = "response")
pred.log.reg <- rep("Down", length(prob.log.reg))
pred.log.reg[prob.log.reg > 0.5] <- "Up"
table(pred.log.reg, Direction)
```

```
##           Direction
## pred.log.reg Down  Up
##           Down   54  48
##           Up    430 557
```

Based on the results of the table above, We may conclude that the percentage of correct predictions (Down \* Down & Up \*Up) on the training data is  $(54 + 557)/1089$  which is equal to 56.11%. So, we can say that 43.89% is the training error rate.

If we look at the data from another angle , meaning we could also conclude that for the *weeks* when the market goes **Up**, the model is right 92.07% of the time  $(557/(48 + 557))$ .

Similarly, for the *weeks* when the market goes **Down**, the model is right only 11.16% of the time  $(54/(54 + 430))$ .

## Question d

- d. Now fit the logistic regression model using a training data period from 1990 to 2008. Compute the confusion matrix and the overall fraction of correct predictions for the held out data (that is, the data from 2009to2010).

## Answer d



*##Part (d) Logistic regression with data from 2009-2010 and the only predictor being "Lag2"*

```
train.data <- (Year < 2009)
Weekly.2009.2010 <- Weekly[!train.data, ]
Direction.2009.2010 <- Direction[!train.data]
log.reg.lag2 <- glm(Direction ~., data = Weekly, family = binomial, subset = train.data)
summary(log.reg.lag2)
```

```
##
## Call:
## glm(formula = Direction ~ ., family = binomial, data = Weekly,
##      subset = train.data)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -1.883e-03 -2.000e-08  2.000e-08  2.000e-08  1.570e-03
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept) -3.258e+03  1.437e+06  -0.002    0.998
## Year         1.632e+00  7.213e+02   0.002    0.998
## Lag1        -4.830e+00  1.233e+03  -0.004    0.997
## Lag2         7.448e+00  1.083e+03   0.007    0.995
## Lag3         1.445e+00  9.872e+02   0.001    0.999
## Lag4         7.540e-01  6.473e+02   0.001    0.999
## Lag5         1.185e+01  1.320e+03   0.009    0.993
## Volume      -8.664e+00  4.189e+03  -0.002    0.998
## Today        8.160e+02  1.686e+04   0.048    0.961
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 1.3547e+03  on 984  degrees of freedom
## Residual deviance: 9.2831e-06  on 976  degrees of freedom
## AIC: 18
##
## Number of Fisher Scoring iterations: 25
```

### ##Part (d) Confusion Matrix

```
prob2.log.reg <- predict(log.reg.lag2, Weekly.2009.2010, type = "response")
pred2.log.reg <- rep("Down", length(prob2.log.reg))
pred2.log.reg[pred2.log.reg > 0.5] <- "Up"
table(pred2.log.reg, Direction.2009.2010)
```

```
##              Direction.2009.2010
## pred2.log.reg Down Up
##              Down  43  0
##              Up   0  61
```

Based on the results of the table above, we can conclude that the percentage of correct predictions on the test data is  $(43 + 61)/104$  (Down \* Down & Up \* Up) which is equal to 100%. So, we can say that the test error rate is 0%.

If we look at the data from another angle , meaning we could also conclude that for the weeks when the market goes **Up** or **Down**, the model is correct 100% of the time.

## Question e

e. Repeat (d) using *LDA*.

## Answer e

### ##Part (e) 1st part - Repeating part (d) using LDA

```
library(MASS)
fit.lda <- lda(Direction ~., data = Weekly, subset = train.data)
fit.lda
```

```
## Call:
## lda(Direction ~ ., data = Weekly, subset = train.data)
##
## Prior probabilities of groups:
##      Down      Up
## 0.4477157 0.5522843
##
## Group means:
##      Year      Lag1      Lag2      Lag3      Lag4      Lag5
## Down 1999.295  0.28944444 -0.03568254 0.17080045 0.15925624 0.21409297
## Up   1998.853 -0.009213235  0.26036581 0.08404044 0.09220956 0.04548897
##      Volume      Today
## Down 1.266966 -1.687018
## Up   1.156529  1.603956
##
## Coefficients of linear discriminants:
##      LD1
## Year   -0.0106936942
## Lag1    0.0003606345
## Lag2    0.0169738374
## Lag3    0.0295058746
## Lag4   -0.0155046298
## Lag5   -0.0279798179
## Volume  0.0587137582
## Today   0.6322256639
```

*##Part (e) 2nd part - Repeating part (d) using LDA*

```
pred.e.lda <- predict(fit.lda, Weekly.2009.2010)
table(pred.e.lda$class, Direction.2009.2010)
```

```
##      Direction.2009.2010
##      Down Up
## Down   36  0
## Up     7 61
```

Based on the results, we conclude that the output is similar but not exactly the same as part (d), which means in this case, the **Logistic Regression** and **LDA** has different results.

From the results of the table above, we can conclude that the percentage of correct predictions on the test data is  $(36 + 61)/104$  (Down \* Down & Up \* Up) which is equal to 93.2%. So, we can say that the test error rate for LDA is 6.8%.

## Question f

f. Repeat (d) using Partial least squares discriminant analysis.

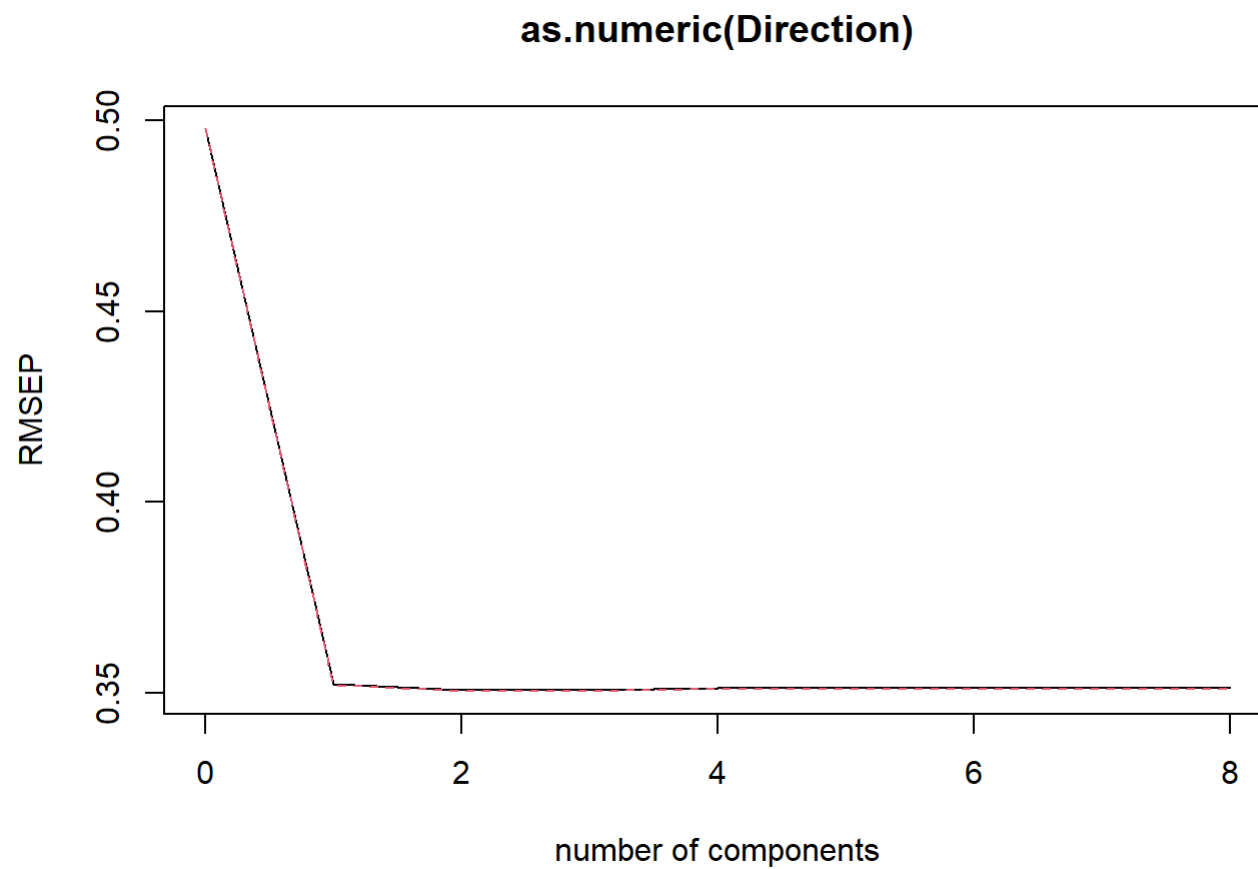
## Answer f

```
# Implementing PLS using plsr() function
```

```
set.seed(1)
fit.pls <- plsr(as.numeric(Direction) ~ ., data = Weekly, subset = train.data, scale=TRUE, validation="CV")
summary(fit.pls)
```

```
## Data:    X dimension: 985 8
## Y dimension: 985 1
## Fit method: kernelpls
## Number of components considered: 8
##
## VALIDATION: RMSEP
## Cross-validated using 10 random segments.
##      (Intercept) 1 comps 2 comps 3 comps 4 comps 5 comps 6 comps
## CV      0.4978   0.3524   0.3507   0.3509   0.3514   0.3513   0.3513
## adjCV    0.4978   0.3522   0.3505   0.3506   0.3510   0.3510   0.3510
##      7 comps 8 comps
## CV      0.3513   0.3513
## adjCV    0.3510   0.3510
##
## TRAINING: % variance explained
##      1 comps 2 comps 3 comps 4 comps 5 comps 6 comps
## X      13.68   32.67   50.47   62.33   66.89   77.18
## as.numeric(Direction) 50.31   51.76   51.87   51.89   51.90   51.90
##      7 comps 8 comps
## X      88.64   100.0
## as.numeric(Direction) 51.90   51.9
```

```
validationplot(fit.pls ,val.type="RMSEP")
```



```
pred.e.pls <- predict(fit.pls, Weekly.2009.2010, ncomp=2)
table(pred.e.pls, Direction.2009.2010)
```

```
##                               Direction.2009.2010
## pred.e.pls                   Down Up
## 0.416864761405766           1 0
## 0.568827594842175           1 0
## 0.582271494839188           1 0
## 0.739425503630556           1 0
## 0.821620099806649           1 0
## 0.831697873207287           1 0
## 0.839805488936676           1 0
## 0.840489933167065           1 0
## 0.887289870391909           1 0
## 0.907048147711969           1 0
## 0.930789176064328           1 0
## 0.945331025137073           1 0
## 0.983787304279044           1 0
## 1.00864454452195            1 0
## 1.12318347594407            1 0
## 1.14920098387987            1 0
## 1.18067226652948            1 0
## 1.1863541979638             1 0
## 1.19655938436577            1 0
## 1.23277276087046            1 0
## 1.25800281085742            1 0
## 1.26067394745872            1 0
## 1.26816783050745            1 0
## 1.3067579725455             1 0
## 1.31983876173129            1 0
## 1.35216167640392            1 0
## 1.36829546833293            1 0
## 1.39057881671399            1 0
## 1.41547156327689            1 0
## 1.41686643434701            1 0
## 1.42279675023241            1 0
## 1.42448310380898            1 0
## 1.43409433932698            1 0
## 1.44131727135222            1 0
## 1.45474035742736            1 0
## 1.45726489400392            1 0
## 1.48661542814847            1 0
```

##	1.494985923391	1	0
##	1.50550722147485	1	0
##	1.51065189513061	1	0
##	1.51209376287901	1	0
##	1.51250935956369	0	1
##	1.51400734316725	1	0
##	1.5440980027812	0	1
##	1.54964442417551	0	1
##	1.55322120103711	0	1
##	1.57603053362911	0	1
##	1.59784838913319	1	0
##	1.60056921924452	0	1
##	1.60918073276943	0	1
##	1.61109474454666	0	1
##	1.62021461031825	0	1
##	1.62212442905777	0	1
##	1.63399155924961	0	1
##	1.65530256822212	0	1
##	1.66324114548606	0	1
##	1.66869175355046	0	1
##	1.68084536990454	0	1
##	1.68409870375885	0	1
##	1.69555854289287	0	1
##	1.69932369701547	0	1
##	1.70817415850165	0	1
##	1.71179460330613	0	1
##	1.71454943060704	0	1
##	1.72269039364966	0	1
##	1.73008506898438	0	1
##	1.73559084886382	0	1
##	1.73887200963136	0	1
##	1.77151490885743	0	1
##	1.79308529877918	0	1
##	1.79987961114314	0	1
##	1.81536913403994	0	1
##	1.82093336190006	0	1
##	1.82940463501608	0	1
##	1.83119892081229	0	1
##	1.8336647256317	0	1



```
## 1.85051947708256      0 1
## 1.8896139081596      0 1
## 1.89003772653626      0 1
## 1.90715430847037      0 1
## 1.91170908820663      0 1
## 1.92005972537925      0 1
## 1.9371947496612      0 1
## 1.96336243568322      0 1
## 1.96772198088562      0 1
## 1.98315072394469      0 1
## 1.98575608964938      0 1
## 1.9892487916503      0 1
## 2.06433026573221      0 1
## 2.06510948331842      0 1
## 2.07344131915426      0 1
## 2.11113356938556      0 1
## 2.11495067723367      0 1
## 2.11729273554395      0 1
## 2.12953756539186      0 1
## 2.16641249040232      0 1
## 2.17121276887075      0 1
## 2.26163807798678      0 1
## 2.32647499717323      0 1
## 2.42191350282431      0 1
## 2.44082687820175      0 1
## 2.63133869858083      0 1
## 2.65594045639992      0 1
## 3.22925164028569      0 1
```

```
mean((pred.e.pls - as.numeric(Weekly.2009.2010$Direction))^2)
```

```
## [1] 0.114054
```

The test error rate for PLS for  $ncomp = 2$  is 11%

## Question g

g. Repeat (d) using Nearest Shrunken Centroids.

Answer g

```
# Implementing Nearest Shrunk Centroids using pamr function

library(pamr)

ctrl <- trainControl(summaryFunction = twoClassSummary, classProbs = TRUE, savePredictions = TRUE)

nscGrid <- data.frame(.threshold = 0:25)
nscTune <- train(x = as.matrix(Weekly[,1:8]),
  y = Weekly$Direction,
  method = "pam",
  preProc = c("center", "scale"),
  tuneGrid = nscGrid,
  metric = "ROC",
  trControl = ctrl)
```

```
## 11111111111111111111111111111111
```


nscTune

```
## Nearest Shrunken Centroids
##
## 1089 samples
##      8 predictor
##      2 classes: 'Down', 'Up'
##
## Pre-processing: centered (8), scaled (8)
## Resampling: Bootstrapped (25 reps)
## Summary of sample sizes: 1089, 1089, 1089, 1089, 1089, 1089, ...
## Resampling results across tuning parameters:
##
## threshold ROC      Sens      Spec
##      0      0.9948327 0.6201356750 0.9998298
##      1      0.9997970 0.5898867640 1.0000000
##      2      1.0000000 0.5630311618 1.0000000
##      3      1.0000000 0.5357010979 1.0000000
##      4      1.0000000 0.5037749125 1.0000000
##      5      1.0000000 0.4651664733 1.0000000
##      6      1.0000000 0.4169004341 1.0000000
##      7      1.0000000 0.3595070091 1.0000000
##      8      1.0000000 0.3025287315 1.0000000
##      9      1.0000000 0.2366655196 1.0000000
##     10      1.0000000 0.1655621794 1.0000000
##     11      1.0000000 0.0943030044 1.0000000
##     12      1.0000000 0.0431982273 1.0000000
##     13      1.0000000 0.0085297598 1.0000000
##     14      0.7200000 0.0006864315 1.0000000
##     15      0.5000000 0.0000000000 1.0000000
##     16      0.5000000 0.0000000000 1.0000000
##     17      0.5000000 0.0000000000 1.0000000
##     18      0.5000000 0.0000000000 1.0000000
##     19      0.5000000 0.0000000000 1.0000000
##     20      0.5000000 0.0000000000 1.0000000
##     21      0.5000000 0.0000000000 1.0000000
##     22      0.5000000 0.0000000000 1.0000000
##     23      0.5000000 0.0000000000 1.0000000
##     24      0.5000000 0.0000000000 1.0000000
##     25      0.5000000 0.0000000000 1.0000000
##
```


```
## ROC was used to select the optimal model using the largest value.
```

```
## The final value used for the model was threshold = 2.
```

```
plot(nscTune)
```



```
#variable importance  
plot(varImp(nscTune, scale =FALSE))
```



```
#Prediction and Error rates
```

```
pred.pamr <- predict(nscTune, Weekly.2009.2010)
table(pred.pamr, Direction.2009.2010)
```

```
##           Direction.2009.2010
## pred.pamr Down Up
##      Down   24  0
##      Up    19 61
```

Based on the results of the table above, We may conclude that the percentage of correct predictions (Down \* Down & Up \* Up) on the training data is  $(24 + 61)/104$  which is equal to 81.73%. So, we can say that 18.27% is the training error rate.

## Question h

h. Which of these methods appears to provide the best results on this data?

## Answer h

Based on the models that we have run and looking at the test error rate data, we can say that **Logistic Regression** is the best model.